

# Lecture 1: Introduction to Machine Learning with Geometry and Topology

Geometry and Topology in Machine Learning Seminar

June 9th, 2025

Lecture 1:  
Introduction  
to Machine  
Learning  
with  
Geometry  
and  
Topology

What is  
Machine  
Learning?

What is a  
Neural  
Network?

Convolutional  
Neural  
Networks

Graph  
Neural  
Networks

- 1 What is Machine Learning?
- 2 What is a Neural Network?
- 3 Convolutional Neural Networks
- 4 Graph Neural Networks

# Programming in the Synthetic World

Lecture 1:  
Introduction to Machine  
Learning  
with  
Geometry  
and  
Topology

What is  
Machine  
Learning?

What is a  
Neural  
Network?


Convolutional  
Neural  
Networks

Graph  
Neural  
Networks

**Q:** Program a function whose (partial) inputs and outputs are:

Input	0	1	3	4.5
Output	5	7	11	14

**A:** Why not consider the function  $f(x) = 2x + 5$ ?



# Learning in the Real World is Hard

Lecture 1:  
Introduction  
to Machine  
Learning  
with  
Geometry  
and  
Topology

What is  
Machine  
Learning?

What is a  
Neural  
Network?

Convolutional  
Neural  
Networks

Graph  
Neural  
Networks

Many complications arise in the real world, including but not limited to:

- ① There may not be an obvious function **explicitly programmed by hand** that **fits the given data**.
- ② Even if such functions exist, they may not **generalize well to additional data**.
- ③ The given data may not have **outputs**, rather you are trying to **learn something out of the data**. The **goal may not be well-defined**.

Rather than trying to construct an explicit program by hand, **machine learning (ML)** tries to find a suitable model from a family of models.

# The General Strategy of Machine Learning

The general strategy of ML tackles these problems in three steps.

- 1 **Function Class:** Define a family of functions (ie. models)  $\mathcal{F}$  that tries to solve the problem given (ex. all functions  $x \mapsto Ax + b$  where  $A \in M_{n,n}(\mathbb{R}), b \in \mathbb{R}^n$ ).  
In practice, a learning problem always reduces to estimating the function  $f$  using a parametrized function class  $\mathcal{F} = \{f_{\theta \in \Theta}\}$
- 2 **Loss/Score Functions:** Define a function  $L : \mathcal{F} \rightarrow \mathbb{R}_{\geq 0}$ <sup>1</sup> that measures error / loss (ex. least square errors).  
Alternatively, this could be a score function for how well the model does.
- 3 **Optimization:** Find the function (or functions)  $f \in \mathcal{F}$  that has the least loss or most score in  $\mathcal{F}$  (ex. Lagrange multipliers).

---

<sup>1</sup>The codomain can be made more general.

# Example: Supervised Learning

**Given:** Data consisting of  $n$  **labeled samples**

$\mathcal{D} = \{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)\}$ , where each  $\mathbf{x}_i \in \mathbb{R}^d$  is a vector of  $d$  features and  $y_i \in \mathbb{R}$  is a label.

**Goal:** Find/Train a model using  $\mathcal{D}$  such that it can predict the label of unknown/future inputs, i.e.,  $\mathbb{P}(Y = y|X = x)$ .

In practice, the inputs will be a **feature matrix** of  $n$  samples where each one  $\mathbf{x}_i$  with  $d$  features, and a label vector  $y \in \mathcal{Y} \subset \mathbb{R}^n$

$$\mathbf{X} = \begin{bmatrix} \text{---}\mathbf{x}_1\text{---} \\ \text{---}\mathbf{x}_2\text{---} \\ \dots \\ \text{---}\mathbf{x}_n\text{---} \end{bmatrix} \quad \text{and} \quad y = \begin{pmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{pmatrix}$$

Here we assume  $\{(\mathbf{x}_i, y_i)\}_{i=1}^n$  **drawn i.i.d. from some unknown distribution**  $\mathbb{P}(x, y)$ .

Given known data  $\mathcal{D} = \{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)\}$  and a hypothesis class  $\mathcal{F}$ :

We can either think of ML as choosing the parameters that minimize the model's average loss on the training data

- **Empirical Risk Minimization (ERM)**: minimize average loss,

$$\theta^* = \arg \min_{\theta \in \mathcal{F}} \frac{1}{n} \sum_{i=1}^n \ell(f_{\theta}(x_i), y_i).$$

In this case, the loss function is

$$\theta \in \mathcal{F} \mapsto \frac{1}{n} \sum_{i=1}^n \ell(f_{\theta}(x_i), y_i),$$

where  $f_{\theta}(x_i)$  is the evaluation of  $\theta$  on  $x_i$  and  $\ell : \mathcal{D} \rightarrow \mathbb{R}_{\geq 0}$  is some function.

Or as finding the parameters that make the observed training data most probable under your model, i.e., maximize the likelihood of prediction

- **Maximum Likelihood Estimation (MLE)**: maximize data likelihood,

$$\theta^* = \arg \max_{\theta \in \mathcal{F}} \prod_{i=1}^n p_{\theta}(y_i | x_i; \theta)$$

In this case, the score function is

$$\theta \in \mathcal{F} \mapsto \prod_{i=1}^n p(y_i | x_i; \theta),$$

where  $\theta \mapsto p(y_i | x_i; \theta)$  is some likelihood function.

Note MLE is equivalent to ERM when  $\ell = -n \log p$ . They together provide two lens of almost all tasks in ML.



In **Part A: Groups and Representations**, our primary focus is on **designing the family of functions  $\mathcal{F}$** .

In this lecture, we will introduce two examples of **neural networks**, which falls in a sub-branch of ML known as **deep learning**:

- 1 Convolutional Neural Networks
- 2 Graph Neural Networks

These two architectures will give us examples for why geometry/topology comes into machine learning.

Lecture 1:  
Introduction  
to Machine  
Learning  
with  
Geometry  
and  
Topology

What is  
Machine  
Learning?

What is a  
Neural  
Network?

Convolutional  
Neural  
Networks

Graph  
Neural  
Networks

- 1 What is Machine Learning?
- 2 What is a Neural Network?**
- 3 Convolutional Neural Networks
- 4 Graph Neural Networks

# What is a Neural Network?

Lecture 1:  
Introduction  
to Machine  
Learning  
with  
Geometry  
and  
Topology

What is  
Machine  
Learning?

What is a  
Neural  
Network?

Convolutional  
Neural  
Networks

Graph  
Neural  
Networks

## Definition

A (feed forward) **neural network** is a **composition of functions**  $L_N \circ L_{N-1} \circ \dots \circ L_1$  where

- 1  $L_i : \mathcal{F}_i \rightarrow \mathcal{F}_{i+1}$  is called the **i-th layer** of the neural network.
- 2  $\mathcal{F}_i$  is called the **i-th feature space**.

- 1 Typically the  $\mathcal{F}_i$ 's are some **real vector spaces**
- 2  $L_0$ , the feature space where the domain lives, is sometimes called the **input layer**. There is no actual function here (or you can think it as an identity map  $L_0(x) = x$ ).  $L_N$  is sometimes called the **output layer**, and  $L_1, \dots, L_{N-1}$  are sometimes called **hidden layers**.

# Example: Perceptrons (Supervised Algorithm)

Lecture 1:  
Introduction  
to Machine  
Learning  
with  
Geometry  
and  
Topology

What is  
Machine  
Learning?

What is a  
Neural  
Network?

Convolutional  
Neural  
Networks

Graph  
Neural  
Networks

## Definition

A **perceptron** is a single layer neural network  $L : \mathbb{R}^n \rightarrow \mathbb{R}^m$  such that, for each  $\mathbf{x} = (x_1, \dots, x_n) \in \mathbb{R}^n$

- ① Given a new data point  $\mathbf{x}_{new}$ , the  $i$ -th component  $L_i(\mathbf{x}_{new}) = \sigma(w_i^\top \mathbf{x}_{new} + b_i)$  computes the activation (or score) for class  $i$ .
- ② where  $w_i \in \mathbb{R}^n$  is some vector (called **weight**) and  $b_i \in \mathbb{R}$  is some real number (called **bias**).
- ③ and  $\sigma$  is some **non-linear** function (called an **activation function**).

Typically,  $\hat{y} = \arg \max_{1 \leq i \leq m} L_i(\mathbf{x}_{new})$  is the predicted label, and we **optimize the choice of  $w_i, b_i$ 's to find the best model**.

An example task of perceptron can do is **binary classification with linear predictor**.

In such case, a perceptron is just the indicator of which side of a hyperplane you're on, i.e,

$$x \mapsto \chi(w^\top x + b > 0).$$

# Example: Multilayer Perceptrons

Lecture 1:  
Introduction to Machine  
Learning with  
Geometry and  
Topology

What is  
Machine  
Learning?

What is a  
Neural  
Network?

Convolutional  
Neural  
Networks

Graph  
Neural  
Networks

## Definition

A **multilayer perceptron** (MLP) is a neural network  $L_N \circ \dots \circ L_1$  such that  $N > 1$  and each  $L_i$  is a perceptron.

Note that the weight matrices and bias vectors at each layer need not be the same.

Here are some examples of common **activation functions/nonlinearities** used:

① **Sign Function**  $\sigma(x) = \begin{cases} 1, x > 0 \\ 0, x = 0 \\ -1, x < 0 \end{cases}.$


② **Hyperbolic tangent functions**

③ **Rectified Linear Units (ReLU)**:  $\sigma(x) = \frac{x+|x|}{2}.$

Note the importance for  $\sigma$  to be **non-linear**. Otherwise, an MLP with linear  $\sigma$ 's is equivalent to a **perceptron** with linear  $\sigma$ .

## MLPs in Image Processing

Consider the following  $28 \times 28$  pixel image<sup>2</sup>  $P$ :




**Question:** Can you classify which number  $N(P)$  that  $P$  is?  
One idea is to solve this using an MLP:

- 1 The input layer is  $\mathbb{R}^{28 \times 28}$ , where each pixel has a **brightness value** between 0.0 (black) and 1.0 (white).
- 2 The output is  $(p(N(P) = 1|P), \dots, p(N(P) = 9|P)) \in \mathbb{R}^9$ .
- 3 This is now a **maximum likelihood estimation** problem.

<sup>2</sup>Picture from the 3Blue1Brown YouTube channel, and MNIST dataset.

Below is a fully connected neural network.



Picture from 3Blue1Brown.

Each neuron has a **basic unit**  $\sigma(w^\top x + b)$ .

If you're seeing this illustration for the first time, there are two questions you may have:

- 1 What do hidden layers represent?
- 2 How are these layers connected?



# Hidden Layers

Lecture 1:  
Introduction to Machine  
Learning with  
Geometry and  
Topology


What is  
Machine  
Learning?

What is a  
Neural  
Network?

Convolutional  
Neural  
Networks

Graph  
Neural  
Networks

Heuristically, the hidden neurons are **used to capture** broken into smaller, recognizable subcomponents (features) as following



Picture from 3Blue1Brown.

But as a hard rule, in a fully-connected net<sup>3</sup> those features are **distributed across many neurons and are not tied to specific pixel locations**, so they cannot capture the desired spatial patterns.

CNNs, as we will cover later, work more like this pattern.


---

<sup>3</sup>ex. MLPs

# Weighted-Sum Mechanism

We want to construct such features using all these neurons.

The key, still **heuristically**, is just a simple word **learnable weighted sum**. Each neuron, or we say each pixel in the first layer, contribute a bit to the output feature of the next layer.



Picture from 3Blue1Brown.

# Universal Approximation

Lecture 1:  
Introduction  
to Machine  
Learning  
with  
Geometry  
and  
Topology

What is  
Machine  
Learning?

What is a  
Neural  
Network?

Convolutional  
Neural  
Networks

Graph  
Neural  
Networks

The choice of perceptrons may seem arbitrary, but they turn out to approximate functions quite effectively! This is known as **universal approximation**.

We state this informally as follows. Let  $K \subset \mathbb{R}^n$  be a nice compact subset<sup>4</sup> and  $f : K \rightarrow \mathbb{R}$  be continuous

- ① The function  $f$  can be arbitrary approximated by a perceptron (but possibly with high width).
- ② Even when  $f$  is not continuous, in many reasonable cases<sup>5</sup>, there is a way to approximate  $f$  using ReLU-based MLPs with minimal width.

---

<sup>4</sup>Think of a shape

<sup>5</sup>For analysts, Bochner–Lebesgue  $p$ -integrable functions

# Drawbacks of MLPs

Lecture 1:  
Introduction to Machine  
Learning  
with  
Geometry  
and  
Topology

What is  
Machine  
Learning?

What is a  
Neural  
Network?

Convolutional  
Neural  
Networks

Graph  
Neural  
Networks

Image classification with MLPs have several limitations:

- 1 If the image has more pixels, the computation costs could go up to be really high.
- 2 There is no spatial structure on the image since it is flattened as a vector (nuisances in the image reorder the inputs arbitrarily). In practice, MLPs do not perform well under [translations of images](#).

In practice, however, many problems involving pictures from data have a natural notion of [translational invariance](#) or [translational equivariance](#)<sup>6</sup> encoded in them:

- 1 What number is the digit in the picture? (Invariance)
- 2 Where is the number in the picture? (Equivariance)

---

<sup>6</sup>The general definition of invariance and equivariance will be given next lecture.

Lecture 1:  
Introduction  
to Machine  
Learning  
with  
Geometry  
and  
Topology

What is  
Machine  
Learning?

What is a  
Neural  
Network?

Convolutional  
Neural  
Networks

Graph  
Neural  
Networks

- 1 What is Machine Learning?
- 2 What is a Neural Network?
- 3 Convolutional Neural Networks**
- 4 Graph Neural Networks

**Observation:** Rather than trying to find a model that is relatively well-behaved with translations, why don't we just restrict our family  $\mathcal{F}$  to neural networks that are **translational invariant or equivariant**?

**Slogan:** "Translation equivariant linear maps are convolutions."

Recall for two functions  $f, g : \mathbb{R} \rightarrow \mathbb{R}$ , their convolution is defined as

$$f * g(t) = \int_{\mathbb{R}} f(\tau)g(t - \tau)d\tau.$$

In the **discrete world**, for two signal  $x, w : \mathbb{Z} \rightarrow \mathbb{R}$ , their convolution is

$$(x * w)_k = \sum_{\tau=-\infty}^{\infty} x_{\tau}w_{k-\tau}, k \in \mathbb{Z}.$$

For two matrices<sup>7</sup>  $x, w : \mathbb{Z}^2 \rightarrow \mathbb{R}$ , their **convolution** is:

$$(x * w)_{i,j} = \sum_{m \in \mathbb{Z}} \sum_{n \in \mathbb{Z}} x(m,n)w(i - m, j - n).$$

---

<sup>7</sup>You can set  $x, w$  to be 0 outside of its grid.

# Convolutions are Translation Equvariant!

Lecture 1:  
Introduction  
to Machine  
Learning  
with  
Geometry  
and  
Topology

What is  
Machine  
Learning?

What is a  
Neural  
Network?

Convolutional  
Neural  
Networks

Graph  
Neural  
Networks

## Lemma:

Convolutions are translation equvariant.

**Proof:** Suppose  $(a, b) \in \mathbb{Z}^2$ , then we observe that

$$\begin{aligned}
 (x * w)_{i-a, j-b} &= \sum_{m \in \mathbb{Z}} \sum_{n \in \mathbb{Z}} x(m, n) w((i-a) - m, (j-b) - n) \\
 &= \sum_{m' = m+a \in \mathbb{Z}} \sum_{n' = n+b \in \mathbb{Z}} x(m' - a, n' - b) w(i - m', j - n') \\
 &= (x_{\bullet-a, \bullet-b} * w)_{i, j}.
 \end{aligned}$$



# Convolutional Neural Network

While there is no set definition for **convolutional neural network** (CNNs), a general paradigm is that it is a neural network that involves a **convolutional layer**.

View a **stack of feature maps** (ex. images) as a map  
 $f : \mathbb{Z}^2 \rightarrow \mathbb{R}^{K_\ell}$ .

## Definition

A **convolutional layer** is a map that takes a stack of feature maps  $f : \mathbb{Z}^2 \rightarrow (\mathbb{R})^{K_\ell}$ :

$$f \mapsto \{(x, y) \rightarrow \left\{ \sum_{(m,n) \in \mathbb{Z}^2} \sum_{i=1}^{K_\ell} f_i(m, n) \psi_i^k(x-m, y-n) \right\}_{k=1, \dots, K_{\ell+1}} \}$$

where  $\psi_i^k : \mathbb{Z}^2 \rightarrow \mathbb{R}$  are a collection of matrices known as **filters**.

Here the filters are not fixed and is learned / optimized by the network.

# Example of a Channel in a Convolutional Layer

Lecture 1:  
Introduction to Machine  
Learning  
with  
Geometry  
and  
Topology

What is  
Machine  
Learning?

What is a  
Neural  
Network?

Convolutional  
Neural  
Networks

Graph  
Neural  
Networks

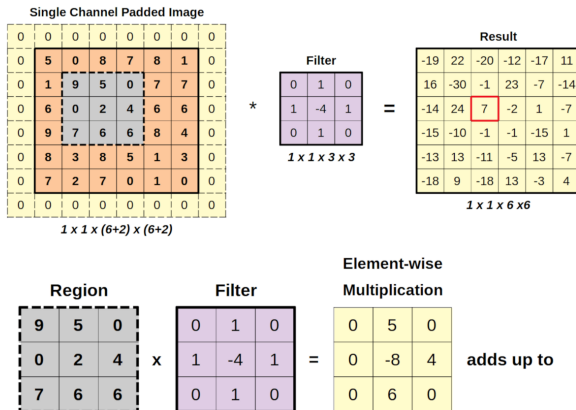


Figure: Visualization by Daniel Godoy.

Although we will not get into details:

- 1 In practice, there is a **stride length** that controls “how far the filters are moving”.
- 2 There is another layer called **pooling** that ensures some translational invariance in the model. The pooling is down typically runs through  $k \times k$  sub-images of the current image and takes the max or average of pixel values.
- 3 A typical CNN could go something in the order of:

Convolution, Pooling, Convolution, Pooling, ...,

Pooling, Fully Connected Layer, Readout.

More generally, one can view the **convolutional layer** of CNN abstract as an **integral transform**:

$$I_{\kappa} : L^2(\mathbb{R}^2, \mathbb{R}^{c_{in}}) \rightarrow L^2(\mathbb{R}^2, \mathbb{R}^{c_{out}})$$

where

- $L^2(\mathbb{R}^2, \mathbb{R}^n)$  is the collection of functions  $f : \mathbb{R}^2 \rightarrow \mathbb{R}^n$  such that

$$\int_{\mathbb{R}^n} ||f||^2 dx < \infty.$$

- $\kappa : \mathbb{R}^2 \times \mathbb{R}^2 \rightarrow \mathbb{R}^{c_{out} \times c_{in}}$  is called the **kernel functions** and

$$I_{\kappa}[f](x) := \int_{\mathbb{R}^2} \kappa(x, y) f(y) dy$$

Here  $\kappa(x, y) f(y)$  is thought of as a matrix multiplication.

# Fun Remark for Analysis Enjoys

Lecture 1:  
Introduction  
to Machine  
Learning  
with  
Geometry  
and  
Topology

What is  
Machine  
Learning?

What is a  
Neural  
Network?

Convolutional  
Neural  
Networks

Graph  
Neural  
Networks

This is just a fun remark for analysis enjoyers, it is totally okay if the following theorem goes over your head.

**Theorem [Dunford-Pettis, see Theorem 1 of [Duits, 2005]]:**

Let  $\mathcal{K} : L^2(X) \rightarrow L^\infty(X)$  be **linear, bounded operator**, then there exists a kernel  $k \in L^1(X, X)$  such that

$$(Kf)(y) = \int_X k(y, x) f(x) d\mu_X(x)$$

with  $f \in L_2(X)$  and  $d\mu_X$  is a **Radon measure** on  $X$ .

To give a glimpse at some of the interactions between CNNs and geometry.

In [Cohen et al., 2018] considered **spherical images** (on  $S^2$ ) and, corresponding, **spherical CNNs** that can process spherical images.

- The primary idea is to replace the translation group with the rotational group  $SO(3)$ .
- Inner product of spherical signals should be:

$$\langle \psi, f \rangle := \int_{S^2} \sum_{k=1}^K \psi_k(x) f_k(x) d\mu_{S^2}$$

where  $d\mu_{S^2}$  is the **surface measure** on  $S^2$ .

Lecture 1:  
Introduction  
to Machine  
Learning  
with  
Geometry  
and  
Topology

What is  
Machine  
Learning?

What is a  
Neural  
Network?

Convolutional  
Neural  
Networks

Graph  
Neural  
Networks

- 1 What is Machine Learning?
- 2 What is a Neural Network?
- 3 Convolutional Neural Networks
- 4 Graph Neural Networks**

- 1 Typically CNNs process images, which can be viewed as a **matrix with pixel values at each coordinate**.
- 2 On another hand, we can reinterpret an  $n \times m$  matrix as an  $n \times m$  **grid graph**<sup>8</sup> with features (ie. pixel values) on each vertex.
- 3 This leads to a more general questions - can we design **neural networks** to handle graphs?
- 4 This is the study of **graph neural networks (GNNs)**.

---

<sup>8</sup>Or grid graph with diagonals



# Motivation for Graph Neural Networks

Lecture 1:  
Introduction  
to Machine  
Learning  
with  
Geometry  
and  
Topology

What is  
Machine  
Learning?

What is a  
Neural  
Network?

Convolutional  
Neural  
Networks

Graph  
Neural  
Networks

At a high level, a **graph neural networks** (GNN) is a model that takes in data represented as **graphs** and perform tasks on graph data.


There are many examples of data in the world that can be represented as graphs.

- **Molecules** can be represented as multi-graphs. The multiplicity of edges between two nodes is the bond type.
- **Recommendation data** can be represented as bipartite graphs between the users and the items.
- **Financial transactions** can be represented as directed graphs where the users are nodes and transactions are edges.

# Graph Neural Networks in Industry

GNNs have found a number of significant industrial applications. On molecular data,

- GNNs can be trained to predict the potency of a molecule. This led to the (re)discovery of [Halicin](#) in [Stokes et al., 2020].



- GNNs showed that [Halicin](#) has strong antibiotic properties.
- In subsequent laboratory tests, the drug successfully treated [Acinetobacter baumannii](#), which is a bacterium resistant to all previously known anti-biotics.

On recommendation data,

- [Alibaba](#) [Zhu et al., 2019] developed a GNN model (called [AliGraph](#)), which is currently used to support product recommendation and personalized search at Alibaba's E-Commerce platform
- [Uber Eats](#) uses the [GraphSAGE](#) model [Hamilton et al., 2017] to provide personalized restaurant recommendations to users.

On financial transactions,

- [Amazon](#) researchers have built a real-time fraud detection system with graph neural networks and the Amazon Web Services (AWS) framework.

Let us look at GNNs. Typically the input of a GNN is the following.

### Definition (Graph)

A **graph** (with features) is the data  $G = (V, E, c, X)$  where  $V$  is a finite vertex set,  $E \subseteq V \times V$ ,  $X$  is a space of features, and  $c : V \rightarrow X$  is a function that assigns each vertex a value in  $X$ .

Let  $v \in G$ , its **neighboring vertices** are denoted  $N_G(v) = \{u \in V : (u, v) \in E\}$ .

Similar to how CNNs look at adjacent pixels in its model, we want to look at **features on adjacent vertices** in designing GNNs.

The **Message-Passing GNN** is typically given as follows,

## Message Passing GNNs

- 1 Initialize  $h_v^{(0)} = c(v)$  for all  $v \in V$ .
- 2 At layer  $t + 1$ , we define the GNN embedding recursively as:

- $m_v^{(t+1)} = \sum_{\omega \in N_G(v)} \text{AGG}^{(t)}(h_v^{(t)}, h_\omega^{(t)})$ .
- $h_v^{(t+1)} = \text{Update}^{(t)}(h_v^{(t)}, m_v^{t+1})$ .

Here  $\text{AGG}^{(t)}$  and  $\text{Update}^{(t)}$  are arbitrary non-linear mappings.

- 3 After reaching the final layer  $T$ , the output is given out using some **readout function**  $\hat{y} = R(\{h_v^{(T)} \mid v \in V\})$ .

Typically the readout function will be used to perform node, edge, or graph classifications.

# Example: Graph Convolutional Networks (GCNs)

Lecture 1:  
Introduction  
to Machine  
Learning  
with  
Geometry  
and  
Topology

What is  
Machine  
Learning?

What is a  
Neural  
Network?

Convolutional  
Neural  
Networks

Graph  
Neural  
Networks

By choosing

$$\text{AGG}^{(t)}(h_v, h_w) = \frac{1}{\sqrt{\deg(v) \deg(w)}} W^{(t)} h_w$$

$$\text{Update}^{(t)}(h_v, m_v) = \sigma(m_v).$$

the general message-passing GNN reduces to special case known as the **Graph Convolutional Network (GCN)**.

Note that the update function on  $v$  does not care about what is happening on the current node.

# Why Do We Need Topology in GNNs?

Lecture 1:  
Introduction to Machine  
Learning with  
Geometry and  
Topology



What is  
Machine  
Learning?

What is a  
Neural  
Network?

Convolutional  
Neural  
Networks

Graph  
Neural  
Networks

Let us take a look at the following 2 graphs.



**Claim:** Our GNN model, at the present, produces the same output on  $G$  and  $H$ .

# Why Do We Need Topology in GNNs?


Lecture 1:  
Introduction to Machine  
Learning with  
Geometry and  
Topology

What is  
Machine  
Learning?


What is a  
Neural  
Network?

Convolutional  
Neural  
Networks

Graph  
Neural  
Networks



$G$



$H$

Recall the aggregate function is given by:

$$\text{AGG}^{(t)}(\{h_w^{(t)} \mid w \in N_G(v)\}).$$

There is no difference in the input to the aggregate function!



Thus, we see that message passing GNNs cannot count **the number of connected components** or **the number of (independent) cycles on a graph**.

Although this will not be the focus of this seminar, if the reader knows some **algebraic topology**, these two data corresponds exactly to:

- 1 The zeroth homology  $H_0(G)$ .
- 2 The first homology  $H_1(G)$ .

We can add **homology** as an additional parameter in the model to increase **expressivity**!

# Why Do We Need Geometry in GNNs?

Lecture 1:  
Introduction  
to Machine  
Learning  
with  
Geometry  
and  
Topology

What is  
Machine  
Learning?

What is a  
Neural  
Network?

Convolutional  
Neural  
Networks

Graph  
Neural  
Networks

Molecules can be viewed as a graph, and **molecules** have a natural embedding in  $\mathbb{R}^3$ :

- 1 This means that one feature each vertex  $v$  of a molecular graph  $G$  can have is a **coordinate**  $x(v) \in \mathbb{R}^3$ .
- 2 We want GNNs to be **equivariant/invariant** on these molecules with respect to translations, rotations, etc. in  $\mathbb{R}^3$ .

More generally, the feature space of a graph could be  $\mathbb{R}^n$  or **any arbitrary vector space  $V$  with the action of some group  $G$** .<sup>9</sup>

More on this next lecture!

---

<sup>9</sup>Also known as a representation.



Cohen, T. S., Geiger, M., Köhler, J., and Welling, M. (2018).

Spherical CNNs.

In *International Conference on Learning Representations*.



Duits, R. (2005).

*Perceptual organization in image analysis : a mathematical approach based on scale, orientation and curvature.*

Phd thesis 1 (research tu/e / graduation tu/e), Biomedical Engineering.



Hamilton, W. L., Ying, R., and Leskovec, J. (2017).

Inductive representation learning on large graphs.

In *Proceedings of the 31st International Conference on Neural Information Processing Systems, NIPS'17*, page 1025–1035, Red Hook, NY, USA. Curran Associates Inc.



Stokes, J. M., Yang, K., Swanson, K., Jin, W., Cubillos-Ruiz, A., Donghia, N. M., MacNair, C. R., French, S., Carfrae, L. A., Bloom-Ackermann, Z., Tran, V. M., Chiappino-Pepe, A., Badran, A. H., Andrews, I. W., Chory, E. J., Church, G. M., Brown, E. D., Jaakkola, T. S., Barzilay, R., and Collins, J. J. (2020).

A deep learning approach to antibiotic discovery.  
*Cell*, 180(4):688–702.e13.



Zhu, R., Zhao, K., Yang, H., Lin, W., Zhou, C., Ai, B., Li, Y., and Zhou, J. (2019).

Aligraph: a comprehensive graph neural network platform.  
*Proc. VLDB Endow.*, 12(12):2094–2105.